



Applied Cryptography

Summer 2026

Part 1: Provable Security

1.3: Provable Security & Computational Cryptography


Nadim Kobeissi

<https://appliedcryptography.page>

Section 1


Provable Security

Last time, we defined subroutines



Subroutines

- “Victim” chooses their key.
- Adversary chooses the message and receives the ciphertext.
- We say that **the adversary has access to an encryption oracle**.



Source: The Joy of Cryptography

Attack scenarios as libraries

- This is a **library** with two **subroutines** and a **global variable** D .
- Victim holds a 6-sided die.
- The first line of the library represents an initialization step.
- Attacker can call the subroutines at any time, and:
 1. Make a guess about the current value of the die and learn whether the guess was correct.
 2. Instruct the victim to (privately) re-roll the die.

$\mathcal{L}_{\text{dice-guess}}$

$D \leftarrow \{1, 2, 3, 4, 5, 6\}$

GUESS(G):

return $D \stackrel{?}{=} G$

REROLL():

$D \leftarrow \{1, 2, 3, 4, 5, 6\}$

One-time pad

From the adversary's perspective...

ATTACK(M):

$K \leftarrow \{\mathbf{0}, \mathbf{1}\}^n$

$C := K \oplus M$

return C

$\approx$

(indistinguishable
from)

JUNK(M):

$C \leftarrow \{\mathbf{0}, \mathbf{1}\}^n$

return C

“Real or random?”

Programs and libraries

- $\mathcal{A}$ is a **program** that calls library $\mathcal{L}_{\text{dice-guess}}$
- Programs can only see the output of function calls to libraries.
 - Programs can't read the values of library variables,
 - Programs can't measure how long it took to run a subroutine,
 - Etc.


$\mathcal{A}$
if GUESS(6): return true REROLL() return GUESS(6)



$\mathcal{L}_{\text{dice-guess}}$
$D \leftarrow \{1, 2, 3, 4, 5, 6\}$ <u>GUESS(G):</u> return $D \stackrel{?}{=} G$ <u>REROLL():</u> $D \leftarrow \{1, 2, 3, 4, 5, 6\}$


When are libraries interchangeable?

- Libraries are **interchangeable** when they:
 - Have the same interface,
 - $\Pr[\mathcal{A} \diamond \mathcal{L}_1 \Rightarrow \text{true}] = \Pr[\mathcal{A} \diamond \mathcal{L}_2 \Rightarrow \text{true}]$
- i.e. when their usage and output is indistinguishable from the adversary's perspective when they are paired with $\mathcal{A}$.
- A lot of the time, $\mathcal{A}$'s mission is to try to *distinguish* between $\mathcal{L}_1$ and $\mathcal{L}_2$.



When are libraries interchangeable?

- Are these libraries interchangeable?
- **Yes!** Their only difference happens in *unreachable lines of code*.




When are libraries interchangeable?

- Are these libraries interchangeable?
- **Yes!** Their only difference is the value they assign to a *variable that is never actually used*.

$\mathcal{L}_a$		$\mathcal{L}_b$
$\frac{\text{FOO}(A, B):}{Y \leftarrow \{\mathbf{0}, \mathbf{1}\}^n}$ $C := \text{BAR}(A)$ return Y	$\stackrel{?}{\equiv}$	$\frac{\text{FOO}(A, B):}{Y \leftarrow \{\mathbf{0}, \mathbf{1}\}^n}$ $C := \text{BAR}(B)$ return Y
$\frac{\text{BAR}(M):}{X \leftarrow \{\mathbf{0}, \mathbf{1}\}^n}$ return X		$\frac{\text{BAR}(M):}{X \leftarrow \{\mathbf{0}, \mathbf{1}\}^n}$ return X


When are libraries interchangeable?

- Are these libraries interchangeable?
- $\parallel$ denotes string concatenation.
 - $\text{Po}\parallel\text{tato} = \text{Potato}$
- **Yes!** Outputting the concatenation of two randomly sampled uniform strings of lengths n and m is the same as outputting a single random string of length $n + m$.



When are libraries interchangeable?

- Are these libraries interchangeable?
- **No!** The first library uses the same K for all subsequent encryptions, which breaks OTP security.
 - Recall that OTP requires keys to be only used once (hence the name).



How do we show the distinguisher?

- What happens if we call $\mathcal{A} \diamond \mathcal{L}_a$ multiple times?
 - We'll get the same output $\equiv K$ each time!
- $\Pr[\mathcal{A} \diamond \mathcal{L}_a \Rightarrow \text{true}] = 1,$


$\mathcal{A}$
$C_1 := \text{OTP.ENC}(\textcolor{red}{0}^n)$ $C_2 := \text{OTP.ENC}(\textcolor{red}{0}^n)$ return $C_1 \stackrel{?}{=} C_2$

$\diamond$

$\mathcal{L}_a$
$K \leftarrow \{\textcolor{red}{0}, \textcolor{red}{1}\}^n$ $\text{OTP.ENC}(M):$ $C := K \oplus M$ return C

How do we show the distinguisher?

- What happens if we call $\mathcal{A} \diamond \mathcal{L}_a$ multiple times?
 - We'll get the same output $\equiv K$ each time!
- $\Pr[\mathcal{A} \diamond \mathcal{L}_a \Rightarrow \text{true}] = 1,$
- $\Pr[\mathcal{A} \diamond \mathcal{L}_b \Rightarrow \text{true}] = \frac{1}{2^n}.$



OTP: why key re-use is bad

$$\begin{aligned}C_i \oplus C_j &= (K \oplus M_i) \oplus (K \oplus M_j) \\&= K \oplus K \oplus M_i \oplus M_j \\&= \mathbf{0}^n \oplus M_i \oplus M_j \\&= M_i \oplus M_j\end{aligned}$$

- $\Pr[\mathcal{A} \diamond \mathcal{L}_a \Rightarrow \text{true}] = 1^a$

$\mathcal{A}$
$M_1 \leftarrow \{\mathbf{0}, \mathbf{1}\}^n$
$M_2 \leftarrow \{\mathbf{0}, \mathbf{1}\}^n$
$C_1 := \text{OTP.ENC}(M_1)$
$C_2 := \text{OTP.ENC}(M_2)$
return $C_1 \oplus C_2 \stackrel{?}{=} M_1 \oplus M_2$

$\diamond$

$\mathcal{L}_a$
$K \leftarrow \{\mathbf{0}, \mathbf{1}\}^n$
$\text{OTP.ENC}(M):$
$C := K \oplus M$
return C

^aInteractive demo:

<https://www.douglas.stebila.ca/teaching/visual-one-time-pad/>

Proving two libraries are interchangeable

- Are these libraries interchangeable?
- Let's find out!
- **Goal:** transform $\mathcal{L}_{\text{xor-samp-1}}$ to $\mathcal{L}_{\text{xor-samp-2}}$, proving that each transformation step does not cause any change whatsoever to the library's functionality.

$\mathcal{L}_{\text{xor-samp-1}}$		$\mathcal{L}_{\text{xor-samp-2}}$
$\frac{\text{SAMPLE}(M):}{X \leftarrow \{\mathbf{0}, \mathbf{1}\}^n}$ $Y := X \oplus M$ return (X, Y)	$\stackrel{?}{\equiv}$	$\frac{\text{SAMPLE}(M):}{Y \leftarrow \{\mathbf{0}, \mathbf{1}\}^n}$ $X := Y \oplus M$ return (X, Y)

Proving two libraries are interchangeable

Step 1

- We start at $\mathcal{L}_{\text{xor-samp-1}}$.

$\mathcal{L}_{\text{xor-samp-1}}$

SAMPLE(M):

$X \leftarrow \{0, 1\}^n$

$Y := X \oplus M$

return (X, Y)

Proving two libraries are interchangeable

Step 2

- Let's add a new variable X' .
- Note that $X = X'$:
 - $X' = Y \oplus M = (X \oplus M) \oplus M = X$.

$\mathcal{L}_{\text{hyb-1}}$

SAMPLE(M):

$X \leftarrow \{0, 1\}^n$

$Y := X \oplus M$

$X' := Y \oplus M$

return (X, Y)

Proving two libraries are interchangeable

Step 3

- So, we can return (X', Y) without any change on the library's effect.
 - $X' = Y \oplus M = (X \oplus M) \oplus M = X$.

$\mathcal{L}_{\text{hyb-2}}$

SAMPLE(M):

$X \leftarrow \{0, 1\}^n$

$Y := X \oplus M$


$X' := Y \oplus M$

return (X', Y)

Proving two libraries are interchangeable

Step 4


- The first two lines of $\text{SAMPLE}(M)$ are the same as $\text{OTP.ENC}(M)$, so we link it and use it instead.



Proving two libraries are interchangeable

Step 5

- The first two lines of $\text{SAMPLE}(M)$ are the same as $\text{OTP.ENC}(M)$, so we link it and use it instead.
- Recall that $\mathcal{L}_{\text{otp-real}} \equiv \mathcal{L}_{\text{otp-rand}}!$
- So, we replace $\mathcal{L}_{\text{otp-real}}$ with $\mathcal{L}_{\text{otp-rand}}$.



Proving two libraries are interchangeable

Step 6

- The first two lines of $\text{SAMPLE}(M)$ are the same as $\text{OTP.ENC}(M)$, so we link it and use it instead.
- Recall that $\mathcal{L}_{\text{otp-real}} \equiv \mathcal{L}_{\text{otp-rand}}!$
- So, we replace $\mathcal{L}_{\text{otp-real}}$ with $\mathcal{L}_{\text{otp-rand}}$.
- We inline $\mathcal{L}_{\text{otp-rand}}$ back into our main library.

$\mathcal{L}_{\text{hyb-5}}$
$\text{SAMPLE}(M):$ $Y \leftarrow \{0, 1\}^n$ $X' := Y \oplus M$ return (X', Y)

Proving two libraries are interchangeable

Step 7

- The first two lines of $\text{SAMPLE}(M)$ are the same as $\text{OTP.ENC}(M)$, so we link it and use it instead.
- Recall that $\mathcal{L}_{\text{otp-real}} \equiv \mathcal{L}_{\text{otp-rand}}!$
- So, we replace $\mathcal{L}_{\text{otp-real}}$ with $\mathcal{L}_{\text{otp-rand}}$.
- We inline $\mathcal{L}_{\text{otp-rand}}$ back into our main library.
- Finally, we rename X' to X .

$\mathcal{L}_{\text{hyb-6}}$
$\text{SAMPLE}(M):$ $Y \leftarrow \{0, 1\}^n$ $X := Y \oplus M$ return (X, Y)

Proving two libraries are interchangeable

Result

- Interchangeable!
- Note how we had to show equivalence **at each step**.

$\mathcal{L}_{\text{hyb-6}}$		$\mathcal{L}_{\text{xor-samp-2}}$
$\frac{\text{SAMPLE}(M):}{Y \leftarrow \{\textcolor{red}{0}, \textcolor{red}{1}\}^n}$	$\equiv$	$\frac{\text{SAMPLE}(M):}{Y \leftarrow \{\textcolor{red}{0}, \textcolor{red}{1}\}^n}$
$X := Y \oplus M$		$X := Y \oplus M$
return (X, Y)		return (X, Y)

Cryptographic primitives

- Specific algorithms like OTP are important in cryptography, but OTP is just one instance of an **encryption scheme**.
- In cryptography, useful abstractions like “encryption scheme” are called **primitives**.
- Three things are important when defining a cryptographic primitive:
 1. **Syntax**: The basic raw interface - algorithms, inputs/outputs and their types.
 2. **Correctness**: Basic functionality without adversaries - e.g., “decryption should be the inverse of encryption”.
 3. **Security**: Guarantees that hold in specific attack scenarios with adversaries.

Cryptographic primitives

Σ : a symmetric-key encryption scheme

- $\Sigma.\text{KeyGen}() = K$
 - Input: none
 - Output: key $K \in \Sigma.\mathcal{K}$ (the “key space”).
- $\Sigma.\text{Enc}(K, M) = C$
 - Input: key $K \in \Sigma.\mathcal{K}$, plaintext $M \in \Sigma.\mathcal{M}$ (the “message space”).
 - Output: ciphertext $C \in \Sigma.\mathcal{C}$.
- $\Sigma.\text{Dec}(K, C) = M$
 - Input: key $K \in \Sigma.\mathcal{K}$, ciphertext $C \in \Sigma.\mathcal{C}$ (the “ciphertext space”).
 - Output: plaintext $M \in \Sigma.\mathcal{M}$.

Correctness of a SKE

Correctness for SKE

An SKE scheme Σ is correct if encryption and decryption are inverses, in the following sense:

$$\Pr[\Sigma.\text{Dec}(K, \Sigma.\text{Enc}(K, M)) = M] = 1$$


for all $M \in \Sigma.\mathcal{M}$ and $K \in \Sigma.\mathcal{K}$.

- The definition involves a probability because $\Sigma.\text{Enc}$ may be a randomized algorithm.
- This means that decryption should **always** recover the original message.
- Even if encryption adds randomness, decryption must be deterministic for each key-ciphertext pair.

One-time secrecy of a SKE

One-time Secrecy for SKE

An SKE scheme Σ has one-time secrecy if the following libraries are interchangeable:



An encryption scheme has one-time secrecy if its ciphertexts are uniformly distributed, when keys are sampled uniformly, kept secret, and used for only one encryption, and no matter how the plaintexts are chosen.

Reminder: AND ($\wedge$)

- Let's replace $\oplus$ with $\wedge$. What would happen?
- Output no longer uniform!

A	B	$A \wedge B$
0	0	0
0	1	0
1	0	0
1	1	1

Table: Truth table for AND operation


ATTACK(M):

$K \leftarrow \{0, 1\}^n$


$C := K \wedge M$

return C

Creating a distinguisher program



- Can you write a distinguisher program showing that these libraries are **not** interchangeable?



- $\Pr[\mathcal{A} \diamond \mathcal{L}_{\text{ots-real}} \Rightarrow \text{true}] = 1$
- $\Pr[\mathcal{A} \diamond \mathcal{L}_{\text{ots-rand}} \Rightarrow \text{true}] = \frac{1}{2^n}$

Section 2

Computational Cryptography

Computational notions in cryptography

- Security definitions in theoretical cryptography can be too strict:
 - Even attacks requiring a trillion years of computation.
 - Even attacks with probability lower than winning the lottery 100 times.
- Modern provable security takes a more practical approach:
 - We dismiss attacks with “ridiculously” high computational cost.
 - We dismiss attacks with “ridiculously” tiny success probability.
- Instead of “no attack can succeed, not even in principle”,
- We prove: “every attack has either ridiculously high cost or ridiculously small success probability”.

The concrete approach to provable security

- In the concrete approach to provable security, we aim to be as quantitative as possible about security claims.
- We rarely say definitively that a cryptographic algorithm “is secure”, as we did with OTP.
- Instead, we use statements like:
 - *“Any attack that expends at most 2^{80} effort can succeed with probability no better than 2^{-64} .”*
- It is up to the user to judge whether this quantitative level of security is acceptable based on their use-case.
- This gives users a concrete basis for security decisions.

Section 2

Computational Cryptography

Subsection 2.1

Asymptotic approach to security

Monetary cost of huge computations


- One way to think about huge computations is their monetary cost.
- The table below shows roughly how much a computation involving 2^n CPU cycles would cost on the cheapest available Amazon EC2 cloud computing service:

Clock cycles	Approx cost (USD)	Point of reference
2^{50}	\$3.50	cup of coffee
2^{55}	\$100	dinner at a high-end restaurant
2^{65}	\$130,000	apartment in Achrafieh
2^{75}	\$130 million	budget of one of the Harry Potter movies
2^{85}	\$140 billion	GDP of Hungary
2^{92}	\$20 trillion	GDP of the United States
2^{99}	\$2 quadrillion	all of human economic activity since 300,000 BCE
2^{128}	a lot!	a billion human civilizations' worth of effort

Exponential vs polynomial growth

Quick side-note

- Important insight: $2^{n+1} = 2 \times 2^n$
 - $2^{51} = 2 \times 2^{50}$
 - $2^{128} = 2^{64} \times 2^{64}$
- Each additional bit doubles the computational cost!
- Compare growth rates:
 - Linear: $f(n) = n$
 - Polynomial: $f(n) = n^2$
 - Exponential: $f(n) = 2^n$



Note: Exponential growth quickly surpasses all polynomial functions!

Understanding tiny probabilities

- Let's put extremely small probabilities in perspective:
 - In 2009, Patricia Demauro rolled dice 154 consecutive times without getting a 7 in craps.
 - Probability of this event: $(30/36)^{154} \approx 2^{-40.5}$
 - Often cited as one of the most improbable documented events in gambling history.
- Other extremely unlikely gambling events:
 - In 1943, a roulette wheel reportedly landed on red 32 consecutive times.
 - Probability: $(18/38)^{32} \approx 2^{-34.5}$
 - Winning the American Powerball lottery: $2^{-28.1}$
 - Winning Powerball two consecutive weeks: $2^{-56.2}$
- These examples help us intuitively grasp what probabilities like 2^{-40} or 2^{-50} actually mean.

Why 2^{-x} for probabilities?

Quick side-note

- We've been using notation like 2^{-40} , 2^{-64} , 2^{-128} for probabilities. Why the negative exponent?
- Recall that $2^{-x} = \frac{1}{2^x}$
 - $2^{-40} = \frac{1}{2^{40}} = \frac{1}{1,099,511,627,776}$
 - This is approximately 1 in a trillion
- Probabilities must be between 0 and 1:
 - 2^{40} would be a huge number (over a trillion) - not a valid probability!
 - 2^{-40} is a tiny fraction - exactly what we want for unlikely events

Computational scale of Bitcoin mining

- Let's consider cryptocurrencies based on proof-of-work, like Bitcoin.
- These systems incentivize users to perform truly obscene amounts of computation.
- In Bitcoin's proof-of-work mechanism:
 - Users race to perform SHA-256 hash computations as fast as possible.
 - The collective Bitcoin network has performed approximately 2^{95} SHA-256 hashes in total.
 - In the last 12 months alone: approximately $2^{93.6}$ hashes.
 - Current market cap for Bitcoin: approximately 1.7 trillion USD.
- This shows the enormous scale of computation that economic incentives can support.

The asymptotic approach to provable security

- The concrete approach gives practical guidance (e.g., key sizes) but requires managing tedious quantitative details.
- The asymptotic approach^a:
 - Makes qualitative, all-or-nothing statements
 - Considers behavior as key size approaches infinity
 - Hides tedious quantitative details
- Similar to asymptotic analysis using big-O notation.
- Example: Instead of calculating that an operation takes exactly $16n^2 + 24n + 74$ steps, we can simply say it takes $\mathcal{O}(n^2)$ steps.

^aAsymptotic means “approaching a limit”. In computer science, we study how algorithms behave as input size approaches infinity, focusing on the general trend rather than exact values.

AES: An example of asymptotic security analysis

- The Advanced Encryption Standard (AES) is a common symmetric-key encryption algorithm.
- It comes in different variants: AES-128, AES-192, AES-256.
- In asymptotic analysis, we'd say:
 - AES is secure against all polynomial-time adversaries.
 - Any successful attack must take exponential time in the key length.
- This hides the concrete details about specific computational costs.
- Concrete security for AES-128:
 - Best known attack: $\approx 2^{126.1}$ operations
- For AES-256:
 - Best known attack: $\approx 2^{254.4}$ operations
 - Far beyond capability of any conceivable computation
- Asymptotic approach: Both are “computationally secure”

Why does a 128-bit key mean 2^{128} possible keys?

- Each bit can be either 0 or 1
- For AES-128:
 - Key has 128 bits
 - Each bit: 2 choices
 - Total combinations: $\underbrace{2 \times 2 \times 2 \times \dots \times 2}_{128 \text{ times}} = 2^{128}$
- That's approximately 3.4×10^{38} different keys!

Polynomial running time

- In the asymptotic approach, we focus only on adversaries that run in polynomial time.
- The security parameter λ determines the level of security:
 - Usually the length of secret keys in bits
 - All algorithms have access to λ as a global variable
 - Example: AES-128 uses $\lambda = 128$, AES-256 uses $\lambda = 256$. ("256-bit security")

Polynomial Running Time

An algorithm runs in polynomial time if there is a polynomial p such that the algorithm takes at most $O(p(n))$ steps on inputs of length n .

Negligible functions

Negligible Functions

A function f is negligible if it approaches zero faster than $1/p(\lambda)$, for every polynomial p . Formally:

- For every polynomial p , there exists λ_0 such that $f(\lambda) < 1/p(\lambda)$ for all $\lambda > \lambda_0$.
- For every polynomial p , we have $p(\lambda) \in O(1/f(\lambda))$.

- You can probably ignore the definition and just think about this intuitively:
- In cryptography, we want attack probabilities to be negligible in the security parameter λ .
- Negligible probability: essentially zero for practical purposes when λ is large enough.

Indistinguishability in computational security

Indistinguishability

Let $\mathcal{L}_1$ and $\mathcal{L}_2$ be two libraries with the same interface. The **advantage** of a calling program $\mathcal{A}$ in distinguishing $\mathcal{L}_1$ and $\mathcal{L}_2$ is:

$$|\Pr[\mathcal{A} \diamond \mathcal{L}_1 \Rightarrow \text{true}] - \Pr[\mathcal{A} \diamond \mathcal{L}_2 \Rightarrow \text{true}]|$$

$\mathcal{L}_1 \approx \mathcal{L}_2$, if every polynomial-time calling program has only negligible advantage in distinguishing them.

- Previously, libraries were interchangeable when probabilities were **identical**.
- Now, libraries are indistinguishable when probabilities are **negligibly close**.
- Again: this is intuitive, no need to stress about the formal definition.

Section 2

Computational Cryptography

Subsection 2.2

Birthday paradox

Birthday paradox

- In a classroom of 50 students, what's the probability that at least two share a birthday?
 - 2%?
 - 20%?
 - 50%?
 - 97%?
- Many people guess around 15-20%, but the actual probability is about 97%!
- This is counterintuitive because **we're not looking for a specific birthday - we're looking for *any* match among all possible pairs.**
- With 50 students, we have $\binom{50}{2} = 1,225$ possible pairs to check!^a

^a $\binom{50}{2}$ is a binomial coefficient. It means: "how many ways can I choose two different items from a set of 50?"

Birthday paradox

- With just 23 people, probability exceeds 50%
- Formula for n people:

$$P = 1 - \frac{365!}{(365 - n)! \cdot 365^n}$$

- Implication: Finding collisions in a space of size N happens with roughly $\sqrt{N}$ samples.
- **This is why many cryptographic systems need large output spaces!**

Birthday probabilities

- Many cryptographic algorithms fail if two executions sample the same random value.
- General question: If we take q independent uniform samples from a set of N items, what's the probability some value is chosen more than once?
- This probability is called $\text{BIRTHDAY}(q, N)$

$$\text{BIRTHDAY}(q, N) = 1 - \prod_{i=1}^{q-1} \left(1 - \frac{i}{N}\right)$$

- Surprisingly high probability!
- With $q \approx 1.2\sqrt{N}$, probability ≈ 0.5
- For birthdays: only need 23 people for $> 50\%$ chance

Section 2

Computational Cryptography

Subsection 2.3

“Bad event” proof technique

The “bad event” proof technique

Definition

Bad Event Technique

Let $\mathcal{L}_1$ and $\mathcal{L}_2$ be libraries that each include a boolean variable named **bad**, and assume that after **bad** is set to **true** it remains **true** forever. We say that the bad event is triggered if the library ever sets **bad** $:=$ **true**.

If $\mathcal{L}_1$ and $\mathcal{L}_2$ have identical source code, except for statements reachable only when **bad** $:=$ **true**, then:

$$|\Pr[\mathcal{A} \diamond \mathcal{L}_1 \Rightarrow \text{true}] - \Pr[\mathcal{A} \diamond \mathcal{L}_2 \Rightarrow \text{true}]| \leq \Pr[\mathcal{A} \diamond \mathcal{L}_1 \text{TRIGGER}(\text{bad})]$$

The “bad event” proof technique


Key insights

- $\mathcal{A}$'s advantage is bounded by $\Pr[\mathcal{A} \diamond \mathcal{L}_1\text{TRIGGER}(\text{bad})]$.
- Practical application:
 - Define a sequence of hybrid libraries.
 - Identify “bad events” between consecutive hybrids.
 - Show these events occur with negligible probability.
- Enables us to focus on analyzing specific failure cases rather than full behavior.

The “bad event” proof technique

Example


- Are these libraries **computationally indistinguishable**?
- Yes! Let’s prove it using the bad event technique.



The “bad event” proof technique


Example

- Are these libraries **computationally indistinguishable**?
- Yes! Let’s prove it using the bad event technique.
- Note: these libraries are interchangeable with $\mathcal{L}_1$ and $\mathcal{L}_2$ (as seen on next slide).



The “bad event” proof technique


Just to be clear



The “bad event” proof technique

Example

- We need to analyze how likely it is that the bad event happens:
 - The bad event occurs when $R \stackrel{?}{=} X$.
 - Since R is randomly chosen from a huge set, this match is extremely unlikely.
 - Even if an adversary makes many attempts, the chance of seeing this bad event remains tiny.
 - As we increase the security parameter λ , the chance becomes vanishingly small.
- Therefore, the libraries are computationally indistinguishable!
- For any $\mathcal{A}$, computationally speaking, we'd get the same distribution of outputs.



The “end-of-time” strategy for bad events

- Sometimes analyzing bad events can be complex, especially when values are chosen by the adversary.
- The end-of-time strategy:
 1. Postpone all bad-event logic to the end of the library execution.
 2. Collect information during normal execution.
 3. Check for bad events only at the very end.
- Advantages:
 - Simplifies analysis by separating normal behavior from bad-event checking.
 - Makes it easier to bound the probability of bad events.
 - Particularly useful for complex cryptographic proofs.
- We'll use this in the next topic!



Applied Cryptography

Summer 2026

Part 1: Provable Security

1.3: Provable Security & Computational Cryptography

Nadim Kobeissi

<https://appliedcryptography.page>