



# Applied Cryptography

Summer 2026


Part 1: Provable Security

## 1.7: Hard Problems & Diffie-Hellman

Nadim Kobeissi

<https://appliedcryptography.page>

# How it's made



Fischer et al, The Challenges of Bringing Cryptography from Research Papers to Products: Results from an Interview Study  
with Experts, USENIX Security 2024

# Cryptographic building blocks

## Security goals

- **Confidentiality:** Data exchanged between Client and Server is only known to those parties.
- **Authentication:** If Server receives data from Client, then Client sent it to Server.
- **Integrity:** If Server modifies data owned by Client, Client can find out.

## Examples

- **Confidentiality:** When you send a private message on Signal, only you and the recipient can read the content.
- **Authentication:** When you receive an email from your boss, you can verify it actually came from them.
- **Integrity:** Your computer can verify that software update downloads haven't been tampered with during transmission.

# Security goals: more examples

- **TLS (HTTPS)** ensures that data exchanged between the client and the server is confidential and that parties are authenticated.
  - Allows you to log into gmail.com without your ISP learning your password.
- **FileVault 2** ensures data confidentiality and integrity on your MacBook.
  - Prevents thieves from accessing your data if your MacBook is stolen.
- **Signal** implements post-compromise security, an advanced security goal.
  - Allows a conversation to “heal” in the event of a temporary key compromise.
  - More on that later in the course.

# Why bother?

- Can't we just use access control?
- Strictly speaking, usernames and passwords can be implemented without cryptography...
- Server checks if the password matches, or if the IP address matches, etc. before granting access.
- What's so bad about that?

## The problem with traditional access control

- Requires trusting the server completely
- No protection during transmission
- No way to verify integrity
- No way to establish trust between strangers

# The magic of cryptography

**Cryptography lets us achieve what seems impossible**

- Secure communication over insecure channels
- Verification without revealing secrets
- Proof of computation without redoing it

Section 1


# Hard Problems

# Hard problems

- Cryptography is largely about equating the security of a system to the difficulty of solving a math problem that is thought to be computationally very expensive.
- With cryptography, we get security systems that we can literally mathematically prove as secure (under assumptions).
- Also, this allows for actual magic.
  - Alice and Bob meet for the first time in the same room as you.
  - You are listening to everything they are saying.
  - Can they exchange a secret without you learning it?



# Time for actual magic



# No known feasible computation

- The discrete logarithm problem:
  - Given a finite cyclic group  $G$ , a generator  $g \in G$ , and an element  $h \in G$ , find the integer  $x$  such that  $g^x = h$
- In more concrete terms:
  - Let  $p$  be a large prime and let  $g$  be a generator of the multiplicative group  $\mathbb{Z}_p^*$  (all nonzero integers modulo  $p$ ).
  - Given:
    - $g \in \mathbb{Z}_p^*, h \in \mathbb{Z}_p^*$
    - Find  $x \in \{0, 1, \dots, p-2\}$  such that  $g^x \equiv h \pmod{p}$
  - This problem is believed to be computationally hard when  $p$  is large and  $g$  is a primitive root modulo  $p$ .
    - “Believed to be” = we don’t know of any way to do it that doesn’t take forever, unless we have a strong, stable quantum computer (Shor’s algorithm)

# Hard problems

## Asymmetric Primitives

- Diffie-Hellman, RSA, ML-KEM, etc.
- “Asymmetric” because there is a “public key” and a “private key” for each party.
- Algebraic, assume the hardness of mathematical problems (as seen just now).

## Symmetric Primitives

- AES, SHA-2, ChaCha20, HMAC...
- “Symmetric” because there is one secret key.
- Not algebraic but unstructured, relying on their understood resistance to  $n$  years of cryptanalysis.
- Can act as substitutes for assumptions in security proofs!
  - Example: hash function assumed to be a “random oracle”

# Hard problems

- Hard computational problems are the cornerstone of modern cryptography.
- These are problems for which even the best algorithms wouldn't find a solution before the sun burns out.
- They provide the security foundation for cryptographic schemes.
- Without hard problems, most of our encryption systems would collapse.

# The rise of computational complexity theory

## Computational Complexity Theory

Complexity theory provides the mathematical framework to understand what makes problems “hard”.

- In the 1970s, rigorous study of hard problems led to computational complexity theory.
- This field has had dramatic impacts beyond cryptography:
  - **Economics:** Computational complexity of finding Nash equilibria in game theory.
  - **Physics:** Simulating quantum many-body systems with exponential complexity.
  - **Biology:** Protein folding prediction and DNA sequence alignment algorithms.

# Computational problems

## Computational Problem

A question that can be answered by performing a computation.

- **Decision problems:** Questions with “yes” or “no” answers
  - Example: “Is 217 a prime number?”
- **Search problems:** Questions that require finding a specific value
  - Example: “How many instances of the letter ‘i’ appear in *‘incomprehensibilities’*?”

- Computational problems form the foundation of theoretical computer science.
- Different types of problems require different algorithmic approaches.
- The difficulty of solving these problems is central to cryptography.

# Computational hardness

## Computational Hardness


The property of computational problems for which no algorithm exists that can solve the problem in a reasonable amount of time.

- Also called **intractable problems**.
- Hardness is independent of the computing device used.
- All standard computing models are equivalent in terms of what they can compute efficiently.
- **Exception:** Quantum computers for certain problems.

- Hardness is a fundamental concept in computational complexity theory.
- Cryptography deliberately uses hard problems to create security.
- What's “hard” should remain hard regardless of hardware advances.

# Measuring algorithm complexity

- To evaluate computational hardness, we need to measure an algorithm's running time.
- We typically use **asymptotic analysis** to express complexity.
- Common notation:
  - $\mathcal{O}(n)$ : Linear time.
  - $\mathcal{O}(n^2)$ : Quadratic time.
  - $\mathcal{O}(2^n)$ : Exponential time.



Complexity classes growth. Source: Serious Cryptography



# Measuring algorithm complexity

- To evaluate computational hardness, we need to measure an algorithm's running time.
- We typically use **asymptotic analysis** to express complexity.
- Common notation:
  - $\mathcal{O}(n)$ : Linear time.
  - $\mathcal{O}(n^2)$ : Quadratic time.
  - $\mathcal{O}(2^n)$ : Exponential time.
- We care about how the running time grows as the input size increases.
- **Example:** An algorithm that takes  $n^2$  operations for input size  $n$  becomes impractical as  $n$  grows large.

# Categorizing computational hardness

## Easy Problems

- Solvable in polynomial time.
- **Examples:** Sorting, searching.
- Running time:  $\mathcal{O}(n^c)$  for some constant  $c$
- Generally scales reasonably with input size.
- Class P (Polynomial time).

## Hard Problems

- No known polynomial-time solution.
- **Example:** Factoring product of two large primes.
- Running time: Often exponential, e.g.,  $\mathcal{O}(2^n)$
- Becomes impractical quickly as input grows.
- Includes NP-hard, NP-complete classes.

# Hard problems in practice


- Public-key cryptography relies on specific hard problems:
  - RSA: Integer factorization problem.
  - Diffie-Hellman: Discrete logarithm problem.
- Cryptography leverages these problems to maximize security assurance.
- The security of these schemes depends on the continued hardness of these problems.

# Quantum vulnerability of hard problems

- The hard problems we rely on today (factoring, discrete logarithm) are vulnerable to quantum computers.
- Shor's algorithm (1994) can efficiently solve both problems on a sufficiently powerful quantum computer.
- This has motivated the search for “**post-quantum**” hard problems:
  - Lattice-based cryptography (e.g., ML-KEM, formerly CRYSTALS-Kyber).
  - Hash-based cryptography.
  - Code-based cryptography.
  - Multivariate cryptography.
  - Isogeny-based cryptography.
- NIST has standardized post-quantum cryptographic algorithms to replace our vulnerable systems.

# What is NIST?

- **NIST** stands for the National Institute of Standards and Technology.
- It's a U.S. government agency that develops technology standards.
- In cryptography, NIST:
  - Sets security standards used worldwide.
  - Evaluates and approves cryptographic algorithms.
  - Led the standardization of post-quantum cryptography.
- When NIST standardizes an algorithm, it often becomes the global industry standard.



NIST's "Standard Reference Peanut Butter",  
available for only \$1,217 USD!

# Funny things standardized by NIST

- **Standard Reference Peanut Butter:** for calibrating food testing equipment.
- **The “Odor Unit”:** for standardizing measurements of smell intensity in environmental monitoring.
- **The Standard Banana Equivalent Dose (BED):** for comparing radiation exposure levels to the natural radiation in a banana.
- **Toilet Paper Testing:** for measuring strength, absorbency, and softness of toilet paper products.

# Cryptographic algorithms standardized by NIST

- **AES (Advanced Encryption Standard):** Selected in 2001 to replace DES, now the worldwide standard for symmetric encryption.
- **SHA-2 and SHA-3 (Secure Hash Algorithms):** Cryptographic hash functions used for digital signatures and data integrity.
- **DSA and ECDSA:** Digital Signature Algorithms based on the discrete logarithm problem.
- **Triple DES:** An interim standard before AES that enhanced the security of the original DES.
- **ML-KEM and ML-DSA:** Recently standardized post-quantum public-key cryptography and signature schemes.

# Why hard problems matter

- Hard problems **create asymmetry between legitimate users and attackers.**
- Easy in one direction, difficult in the reverse.
- Example: Easy to multiply large primes, hard to factor the product.
- This asymmetry is what enables secure communication!



# What are complexity classes?

## Complexity Class

A group of computational problems that share similar resource requirements (time, memory, etc.).

- **Example:** All problems solvable in  $\mathcal{O}(n^2)$  time form one class.
- Different classes represent different levels of computational difficulty.
- Understanding these classes helps us categorize cryptographic problems.

# TIME complexity classes

- **TIME**( $f(n)$ ) = class of problems solvable in time  $\mathcal{O}(f(n))$
- Examples:
  - **TIME**( $n^2$ ) = problems solvable in  $\mathcal{O}(n^2)$  time
  - **TIME**( $n^3$ ) = problems solvable in  $\mathcal{O}(n^3)$  time
  - **TIME**( $2^n$ ) = problems solvable in  $\mathcal{O}(2^n)$  time
- **Key insight:** If you can solve a problem in  $\mathcal{O}(n^2)$  time, you can also solve it in  $\mathcal{O}(n^3)$  time.
- Therefore: **TIME**( $n^2$ )  $\subseteq$  **TIME**( $n^3$ )  $\subseteq$  **TIME**( $n^4$ )  $\subseteq$  ...

# The class P (Polynomial time)

## Class P

The union of all **TIME**( $n^k$ ) classes for all constants  $k$ .

- $P = \text{TIME}(n) \cup \text{TIME}(n^2) \cup \text{TIME}(n^3) \cup \dots$
  - Contains all problems solvable in polynomial time.
  - Generally considered “efficiently solvable”.
- 
- Most practical algorithms we use daily are in class P.
  - Examples: Sorting, searching, basic arithmetic.
  - Cryptography often relies on problems **not** in P!

# SPACE complexity classes

- Time isn't everything—memory usage matters too!
- A single memory access can be orders of magnitude slower than CPU operations.
- **SPACE**( $f(n)$ ) = class of problems solvable using  $\mathcal{O}(f(n))$  bits of memory
- Examples:
  - **SPACE**( $n$ ) = problems using  $\mathcal{O}(n)$  memory
  - **SPACE**( $n^2$ ) = problems using  $\mathcal{O}(n^2)$  memory
- **PSPACE** = union of all **SPACE**( $n^k$ ) for constants  $k$

# Relationship between TIME and SPACE

- **Key insight:** Any algorithm running in time  $f(n)$  uses at most  $f(n)$  memory.
- Why? You can write at most one bit per time unit.
- Therefore:  $\mathbf{TIME}(f(n)) \subseteq \mathbf{SPACE}(f(n))$
- This gives us:  $P \subseteq PSPACE$
- **Important:** Low memory doesn't guarantee fast execution!
  - Example: Brute-force key search uses little memory but takes forever.

# The class NP (Nondeterministic Polynomial time)

## Class NP

The class of decision problems for which you can **verify** a solution in polynomial time, even if finding the solution is hard.

- **Key insight:** Easy to check, hard to find!
- Given a potential solution, you can run a polynomial-time algorithm to verify if it's correct.
- You don't need to find the solution efficiently—only verify it efficiently.
- **Relationship:**  $P \subseteq NP$  (if you can solve it quickly, you can certainly verify it quickly)

# NP: A cryptographic example

**Problem:** Given a hash value  $H$ , does there exist an input  $M$  such that  $\text{HASH}(M) = H$ ?

- **Finding the solution:** Could take exponential time (brute-force searching through possible inputs)
- **Verifying a candidate solution:** Given a potential input  $M_0$ :
  1. Compute  $\text{HASH}(M_0)$
  2. Check if  $\text{HASH}(M_0) = H$
  3. Return “yes” if they match, “no” otherwise
- This verification runs in polynomial time!
- Therefore, this hash preimage problem is in NP.

# NP: Another cryptographic example

**Problem:** Given plaintext  $P$  and ciphertext  $C$ , does there exist a key  $K$  such that  $C = E(K, P)$ ?

- **Finding the solution:** Could take exponential time (brute-force key search)
- **Verifying a candidate solution:** Given a potential key  $K_0$ :
  1. Compute  $E(K_0, P)$
  2. Check if  $E(K_0, P) = C$
  3. Return “yes” if they match, “no” otherwise
- This verification runs in polynomial time!
- Therefore, this key recovery problem is in NP.



# What's NOT in NP?

- **Known-ciphertext attacks:** You only have  $E(K, P)$  values for random unknown plaintexts  $P$ .
  - How do you verify if candidate key  $K_0$  is correct?
  - You don't know what the plaintexts should be!
  - Can't express this as a decision problem with efficient verification.
- **Proving absence of solutions:** "Does there exist NO solution to this problem?"
  - To verify "no solution exists," you might need to check all possible inputs.
  - If there are exponentially many inputs, this takes exponential time.
  - Therefore, proving non-existence is generally not in NP.

# NP-complete problems

## NP-Complete Problems

The hardest decision problems in the class NP.

- No known polynomial-time algorithms exist for worst-case instances.
  - If any NP-complete problem can be solved efficiently, then **all** problems in NP can be solved efficiently.
- 
- Discovered in the 1970s during the development of complexity theory.
  - **Remarkable discovery:** All NP-complete problems are fundamentally equally hard!
  - Examples: Boolean satisfiability (SAT), traveling salesman problem, graph coloring.

# Why are NP-complete problems equally hard?

- **Key insight:** You can *reduce* any NP-complete problem to any other NP-complete problem.
- **Reduction:** Transform one problem into another in polynomial time.
  - If you can solve problem B efficiently, you can solve problem A efficiently too.
- **Mathematical equivalence:** Different NP-complete problems may look completely different but are fundamentally the same from a computational perspective.
- **Consequence:** Solving any single NP-complete problem efficiently would solve *all* problems in NP efficiently!
  - This would prove that  $P = NP$  (one of the biggest open questions in computer science).

# The remarkable equivalence of NP-complete problems

These problems look completely different...

## Boolean Logic

Can you set variables to make this formula true?

$$(x_1 \vee \neg x_2) \wedge (x_2 \vee x_3) \wedge \dots$$

## Travel Planning

What's the shortest route visiting all cities exactly once?

## Sudoku Puzzles

Can you fill this  $9 \times 9$  grid following the rules?

...but they're computationally identical!

# Concrete examples of equivalent problems

- **3-SAT** (Boolean satisfiability): Given a logical formula, can you set the variables to make it true?
  - Example:  $(x_1 \vee \neg x_2 \vee x_3) \wedge (\neg x_1 \vee x_2 \vee \neg x_3) \wedge \dots$
- **Traveling Salesman Problem**: Given cities and distances, what's the shortest route visiting each city exactly once?
  - Looks like a geometry/optimization problem!
- **Graph Coloring**: Can you color a graph's vertices with  $k$  colors so no adjacent vertices share a color?
  - Looks like a combinatorial puzzle!
- **Subset Sum**: Given a set of integers, is there a subset that sums to exactly  $k$ ?
  - Looks like an arithmetic problem!

# The magic of reductions

## Problem Reduction

A polynomial-time transformation that converts any instance of problem A into an equivalent instance of problem B.

- You can transform **any** Sudoku puzzle into a Boolean logic formula!
  - The Sudoku has a solution  $\Leftrightarrow$  the formula is satisfiable
- You can transform **any** traveling salesman instance into a graph coloring problem!
- You can transform **any** Boolean formula into a subset sum problem!
- These transformations preserve the “yes/no” answer and run in polynomial time.
- **Mind-blowing consequence:** Solve Sudoku efficiently = solve all of theoretical computer science!

# Real-world impact of this equivalence

- **Good news:** Any algorithmic breakthrough on one NP-complete problem immediately applies to thousands of others!
  - Better SAT solvers  $\Rightarrow$  better protein folding, circuit design, AI planning...
- **Sobering reality:** 50+ years of computer science research suggests these problems are fundamentally hard.
  - Despite massive incentives (millions in prize money, practical applications worth billions)
- **Cryptographic relevance:** We rely on NP-complete problems being hard for certain security models.
  - Though most practical cryptography uses different hard problems (factoring, discrete log)
- **Universal truth:** The computational universe has these deep, hidden connections that unite seemingly unrelated problems.

# Fun fact: Nintendo games are NP-hard!

- **Games proven NP-hard<sup>a</sup>:**
  - Super Mario Bros. 1-3, The Lost Levels, Super Mario World
  - Donkey Kong Country 1-3
  - All classic Legend of Zelda games
  - All classic Metroid games
  - All classic Pokémon role-playing games
- **The catch:** “Generalized versions” with arbitrarily large levels.
  - Real Nintendo levels are designed to be solvable by humans.
  - But the **mathematical structure** of these games is inherently complex.
- **Cool insight:** Video games naturally encode complex computational problems!

---

<sup>a</sup><https://appliedcryptography.page/paper/#nintendo-hard>



# The P vs. NP Problem

## The P vs. NP Problem

One of the most important unsolved problems in computer science and mathematics.

- **Question:** Does  $P = NP$ ?
  - **Translation:** Are there problems that are easy to verify but fundamentally hard to solve?
- 
- If you could solve **any** NP-complete problem in polynomial time, then you could solve **all** NP problems in polynomial time.
  - This would mean  $P = NP$ .
  - **Intuition says:** Surely some problems are easy to check but hard to find!
  - **Example:** Brute-force key recovery seems inherently exponential-time...
  - **Reality:** No one has proved this mathematically!

# The million-dollar question

- The **Clay Mathematics Institute** offers \$1,000,000 for solving P vs. NP.
- One of seven “Millennium Prize Problems”.
- Renowned complexity theorist Scott Aaronson called it *“one of the deepest questions that human beings have ever asked”*.
- **To win:** Prove either  $P = NP$  or  $P \neq NP$ .
- Over 50 years of research, no solution yet!

# What if $P = NP$ ?

## The cryptographic apocalypse scenario

- If  $P = NP$ , then **any easily checked solution would also be easy to find.**
- **Symmetric cryptography** would be completely broken:
  - Key recovery becomes polynomial-time.
  - AES, ChaCha20, all symmetric ciphers become useless.
- **Hash functions** would be invertible in polynomial time:
  - Finding preimages becomes easy.
  - Digital signatures, password storage, all broken.
- **All of modern cryptography** would collapse overnight!
- **But also:** We could solve protein folding, optimize supply chains perfectly, solve climate modeling...

# Why we don't panic

- **Overwhelming consensus:** Most complexity theorists believe  $P \neq NP$ .
- **Intuitive reasoning:** Problems that look hard actually **are** hard.
- **The structure of reality:** Easy-to-verify but hard-to-solve problems seem fundamental to the universe.
- **50+ years of evidence:** Despite massive incentives, no polynomial-time algorithms found for NP-complete problems.
- **Current belief:**  $P$  is a strict subset of  $NP$ , with NP-complete problems outside  $P$ .

## The Challenge

- **Proving  $P = NP$ :** Need only one polynomial-time algorithm for one NP-complete problem
- **Proving  $P \neq NP$ :** Must prove no such algorithm can **ever** exist—much harder!

# Why NP-complete problems don't work for cryptography

- **Tempting idea:** Base cryptography on NP-complete problems for provable security!
- **The dream:** Prove that breaking some cipher is NP-hard.
  - Security would be guaranteed as long as  $P \neq NP$ .
- **Reality is disappointing:** NP-complete problems are hard in the **worst case**, not the **average case**
  - The structure that makes them hard can make specific instances easy.
  - Cryptography needs problems that are hard for **random** instances.
- **What we actually use:** Problems that are probably **not** NP-hard.
  - Factoring, discrete logarithm, lattice problems.
  - Believed hard on average, but not proven NP-complete.

# NP-complete vs. NP-hard

## NP-Complete Problems

- Must be decision problems (yes/no answers)
- You can verify solutions in polynomial time
- **Examples:** 3-SAT, graph coloring, subset sum
- The “sweet spot” of hardness

## NP-Hard Problems

- Can be any type of problem (optimization, etc.)
- May not have polynomial-time verification
- **Examples:** Traveling salesman optimization, halting problem
- Can be even harder than NP-complete!

# Average-case vs. worst-case hardness

## Worst-case hardness (NP-complete)

- Some instances of the problem are very hard.
- Other instances might be easy.
- **Example:** 3-SAT has hard instances, but also trivial ones.
- Not suitable for cryptography.

## Average-case hardness (Crypto)

- Random instances are typically hard.
  - **Example:** Getting  $a$  or  $b$  if I know  $(g, g^a \bmod p, g^b \bmod p, g^{ab} \bmod p)$  (Diffie-Hellman).
  - **Example:** Factoring products of large primes (RSA).
- Perfect for cryptographic applications.

### Hard Problems for Cryptography


We need problems where almost every instance is hard, not just the worst ones.

Section 2

Diffie-Hellman



# Time for actual magic



# The key exchange problem

- Alice and Bob want to communicate securely over the internet.
- They've never met before and share no secrets.
- How can they establish a shared secret key for encryption?
- Traditional approach: meet in person, exchange keys physically.
- **Problem:** This doesn't scale for the internet!

## The Challenge

Create a shared secret between two parties who have never communicated before, even when an eavesdropper can see everything they send to each other.

# The magic of Diffie-Hellman

- Whitfield Diffie and Martin Hellman solved this “impossible” problem.
- Their solution (1976) came one year **before** RSA (1977).
- Uses the **discrete logarithm problem** as its foundation.
- Allows two strangers to create a shared secret in public!

**This was the birth of modern cryptography**

# What makes discrete logarithm hard?

- Remember: we need problems that are easy in one direction, hard in reverse.
- **Easy direction:** Given  $g$  and  $x$ , compute  $g^x \bmod p$ 
  - Example:  $2^{10} \bmod 17 = 1024 \bmod 17 = 4$
- **Hard direction:** Given  $g$ ,  $p$ , and  $g^x \bmod p$ , find  $x$ 
  - Example: Given  $g = 2$ ,  $p = 17$ , and result  $= 4$ , find  $x = 10$
- For small numbers, this is easy. For huge numbers (thousands of bits), it's computationally infeasible!

# A simple example

Let's work with small numbers to see the pattern:

- Let  $p = 17$  (a prime) and  $g = 2$  (a base)
- **Computing powers is easy:**
  - $2^1 \bmod 17 = 2$
  - $2^2 \bmod 17 = 4$
  - $2^3 \bmod 17 = 8$
  - $2^4 \bmod 17 = 16$
  - $2^5 \bmod 17 = 15$
- **Finding the exponent is harder:**
  - Given result 15, can you quickly find that the exponent was 5?
  - With small numbers: yes, by trying all possibilities
  - With 2048-bit numbers: practically impossible!

# Mathematical groups: the foundation

## What is a Mathematical Group?

A set of elements with an operation that follows specific rules.

- Think of it as a **mathematical playground** with consistent rules.
  - For cryptography, we use  $\mathbb{Z}_p^*$ : numbers  $\{1, 2, 3, \dots, p-1\}$  with multiplication mod  $p$ .
- 
- **Example:**  $\mathbb{Z}_5^* = \{1, 2, 3, 4\}$  with multiplication mod 5
    - $3 \times 4 = 12 \bmod 5 = 2$
    - $2 \times 3 = 6 \bmod 5 = 1$
  - The “rules” ensure the math works consistently for cryptography.

# Group rules (simplified)

For our cryptographic group  $\mathbb{Z}_p^*$ , these rules always hold:

- **Closure:** Multiplying any two elements gives another element in the group
  - In  $\mathbb{Z}_5^*$ :  $2 \times 3 = 1$  (still in the group!)
- **Identity:** There's a special element (1) that doesn't change others
  - $1 \times 4 = 4, 1 \times 2 = 2$ , etc.
- **Inverses:** Every element has a "partner" that multiplies to 1
  - In  $\mathbb{Z}_5^*$ :  $2 \times 3 = 1$ , so 2 and 3 are inverses
- **Associativity:**  $(a \times b) \times c = a \times (b \times c)$

**Why care?** These rules guarantee that our cryptographic operations will behave predictably!

# Generators: allow exponentiation to target every element in the group

## Generator

An element  $g$  whose powers  $g^1, g^2, g^3, \dots$  produce **every element in the group**.

- In  $\mathbb{Z}_5^*$ , let's try  $g = 2$ :
  - $2^1 \bmod 5 = 2$
  - $2^2 \bmod 5 = 4$
  - $2^3 \bmod 5 = 3$
  - $2^4 \bmod 5 = 1$
- We got  $\{2, 4, 3, 1\}$  - that's all elements! So  $g = 2$  is a generator.
- **Generators are crucial:** They let us express every group element as a power of  $g$ .



# The discrete logarithm problem (DLP)

## Discrete Logarithm Problem

Given  $g$ ,  $p$ , and  $h = g^x \bmod p$ , find the secret exponent  $x$ .

- “**Discrete**” because we work with integers, not real numbers
- “**Logarithm**” because we’re finding the exponent (like  $\log_2(8) = 3$ )
- **Example:** Given  $g = 2$ ,  $p = 17$ ,  $h = 8$ , find  $x$  such that  $2^x \equiv 8 \pmod{17}$ 
  - Answer:  $x = 3$  (since  $2^3 = 8$ )
  - Easy with small numbers, hard with large ones!
- For cryptographic-sized numbers (2048+ bits), no efficient algorithm is known.

# DLP vs. factoring: similarly hard

## Factoring Problem

- Given  $N = p \times q$ , find  $p$  and  $q$
- Used in RSA (1977)
- Well-known, intuitive

## Discrete Logarithm

- Given  $g^x \bmod p$ , find  $x$
- Used in Diffie-Hellman (1976)
- Less intuitive, more mathematical

- **Security equivalence:**  $n$ -bit factoring  $\approx$   $n$ -bit discrete logarithm
- Both are vulnerable to Shor's quantum algorithm
- Both are **not** known to be NP-hard
- Algorithms for both problems share similar techniques

# Diffie-Hellman: the mathematical version

**Setup:** Alice and Bob agree on public values  $g$  (generator) and  $p$  (large prime)

1. **Alice:** Chooses secret  $a$ , computes  $A = g^a \bmod p$ , sends  $A$  to Bob
2. **Bob:** Chooses secret  $b$ , computes  $B = g^b \bmod p$ , sends  $B$  to Alice
3. **Alice:** Computes shared secret  $S = B^a \bmod p = (g^b)^a \bmod p = g^{ab} \bmod p$
4. **Bob:** Computes shared secret  $S = A^b \bmod p = (g^a)^b \bmod p = g^{ab} \bmod p$

**Result:** Alice and Bob both have  $S = g^{ab} \bmod p$  without ever sharing  $a$  or  $b$ !

# Diffie-Hellman example with small numbers

## Setup

- Public parameters:  $p = 13, g = 2$
- Alice picks secret:  $a = 5$
- Bob picks secret:  $b = 7$

## Public Exchange

- Alice computes:  $A = g^a \bmod p = 6$
- Bob computes:  $B = g^b \bmod p = 11$
- Alice sends  $A = 6$  to Bob
- Bob sends  $B = 11$  to Alice

## Shared Secret Computation

- Alice computes:  
 $s = B^a \bmod p = 11^5 \bmod 13$ 
  - $= 161051 \bmod 13 = 7$
- Bob computes:  
 $s = A^b \bmod p = 6^7 \bmod 13$ 
  - $= 279936 \bmod 13 = 7$
- **Shared secret:  $s = 7$**

## Eavesdropper sees only:

- $p = 13, g = 2, A = 6, B = 11$
- (In the real world,  $p$ ,  $a$  and  $b$  are much larger numbers)

# The computational Diffie-Hellman (CDH) problem

## Computational Diffie-Hellman (CDH) Problem

Given  $g^a \bmod p$  and  $g^b \bmod p$ , compute the shared secret  $g^{ab} \bmod p$  without knowing the secret values  $a$  and  $b$ .

- **Motivation:** Even if an eavesdropper captures the public values  $g^a$  and  $g^b$ , they shouldn't be able to determine the shared secret  $g^{ab}$ .
- **Example:** Given  $A = 6$  and  $B = 11$  from our earlier example, can you compute  $S = 7$ ?
  - Without knowing  $a = 5$  and  $b = 7$ , this becomes very difficult!
- **Real-world relevance:** This is exactly what an attacker faces when trying to break Diffie-Hellman.

# CDH vs. DLP: the relationship

- **Key insight:** If you can solve DLP, then you can also solve CDH.
  - Given  $g^a$  and  $g^b$ , use DLP to find  $a$  and  $b$
  - Then compute  $g^{ab}$  directly
- **Mathematical relationship:** DLP is **at least as hard** as CDH.
  - $\text{CDH} \leq \text{DLP}$  (CDH reduces to DLP)
- **Open question:** Is CDH at least as hard as DLP?
  - We don't know if solving CDH allows you to solve DLP!
  - Maybe there's a clever way to compute  $g^{ab}$  without finding  $a$  and  $b$
- **Security assumption:** We assume CDH is hard even if it's easier than DLP.

# The decisional Diffie-Hellman (DDH) problem

## Decisional Diffie-Hellman (DDH) Problem

Given  $g^a \bmod p$ ,  $g^b \bmod p$ , and a value  $X$  that is either:

- $g^{ab} \bmod p$  (the real shared secret), or
- $g^c \bmod p$  for some random  $c$

...determine which one  $X$  is (each choice has probability 1/2).

- **Why do we need this?** Indistinguishability!
  - What if an attacker can compute the first 32 bits of  $g^{ab}$ ?
  - CDH isn't completely broken, but the attacker learned something.
  - This partial information might compromise application security.
- **DDH ensures:** The shared secret  $g^{ab}$  is **indistinguishable** from a random group element.

# DDH vs. CDH: the hierarchy

- **Key relationship:** If you can solve CDH, then you can solve DDH.
  - Given  $(g^a, g^b, X)$ , use CDH to compute  $g^{ab}$
  - Check if  $X = g^{ab}$ ; if yes, then  $X$  is the real shared secret
- **Hardness hierarchy:**  $\text{DDH} \leq \text{CDH} \leq \text{DLP}$ 
  - DDH is fundamentally **easier** than CDH.
  - CDH is (probably) easier than DLP.
- **Surprising fact:** DDH is **not hard** in certain groups!
  - In  $\mathbb{Z}_p^*$ , DDH can be broken using the Legendre symbol (quadratic residuosity).
  - But CDH remains hard in the same group.
- **Solution:** Use elliptic curve groups where DDH is believed hard.



# Why DDH matters in cryptography

- **Indistinguishability:** DDH ensures that shared secrets “look random”.
  - Critical for encryption schemes and key derivation.
  - Prevents attackers from learning partial information.
- **Security proofs:** Many cryptographic protocols prove security under DDH.
  - ElGamal encryption.
  - Cramer-Shoup cryptosystem.
  - Various authenticated key exchange protocols.
- **Real-world impact:** Even though DDH is “weaker” than CDH, it’s one of the most studied and used assumptions.
  - Provides stronger security guarantees for applications.
  - Enables more sophisticated cryptographic constructions.

Section 2

Diffie-Hellman

---

Subsection 2.1

Real-World Diffie-Hellman

# Real-world Diffie-Hellman

- **TLS/HTTPS:** Your browser uses Diffie-Hellman to establish secure connections.
- **Signal:** Uses elliptic-curve Diffie-Hellman for key exchange.
- **SSH:** Secure shell connections use Diffie-Hellman for key agreement.
- **VPNs:** Many VPN protocols rely on Diffie-Hellman for establishing tunnels.

## Modern Diffie-Hellman Variants

- **Elliptic Curve Diffie-Hellman (ECDH):** Same idea, different mathematical group.
- **Post-quantum alternatives:** New key exchange methods for the quantum era.

More on both of the above in future course topics!

# Diffie-Hellman key exchange in practice

## How does this actually work in the real world?

1. **Parameter generation:** Choose secure values for  $p$  and  $g$ 
  - $p$  must be a large prime (2048+ bits)
  - $g$  must be a generator of a large subgroup
2. **Key generation:** Each party picks a random secret
  - Alice picks  $a$  randomly from  $\{1, 2, \dots, p - 2\}$
  - Bob picks  $b$  randomly from  $\{1, 2, \dots, p - 2\}$
3. **Public key computation:** Each party computes their public value
4. **Key exchange:** Public values are sent over the network
5. **Shared secret derivation:** Each party computes the final shared secret

# TLS handshake: Diffie-Hellman in action


When you visit <https://gmail.com>, here's what happens:

1. **Client Hello:** Your browser says "I want to talk securely"
2. **Server Hello:** Gmail's server responds with its certificate and DH parameters
  - Includes  $p$ ,  $g$ , and server's public value  $g^b \bmod p$
3. **Client Key Exchange:** Your browser generates its own secret  $a$  and sends  $g^a \bmod p$
4. **Secret computation:** Both sides compute  $g^{ab} \bmod p$
5. **Key derivation:** The shared secret is used to derive encryption keys
6. **Secure communication:** All further messages are encrypted with these keys

**Result:** Your password is encrypted before leaving your computer!

# Signal's double ratchet: DH everywhere

- **Initial key exchange:** Uses X3DH (Extended Triple DH)
  - Combines **three** DH key exchanges for security.
  - Works even when recipient is offline ("*asynchronous*" protocol).<sup>a</sup>
- **Ongoing communication:** Uses Double Ratchet
  - New DH key exchange for every message!
  - Provides "forward secrecy" and "post-compromise security".
  - If your phone gets compromised today, yesterday's messages remain secure.
  - If your phone recovers from compromise, tomorrow's messages are secure again.



Signal uses DH key exchange dozens, hundreds of times per conversation.

---

<sup>a</sup>Everything on this slide will be covered in much more detail later in the course.

Section 2

Diffie-Hellman

---

Subsection 2.2

Attacks on Diffie-Hellman

# Weak groups in the wild: the Logjam findings

## Not all Diffie-Hellman implementations are created equal

- The Logjam paper (2015)<sup>a</sup> revealed widespread weaknesses in real-world DH deployments
- Researchers scanned millions of servers and found shocking vulnerabilities
- Three major categories of problems:
  - Servers using 512-bit primes for non-export DHE.
  - Composite-order subgroup attacks.
  - Misconfigured DSA groups.
- **Observation:** Even correct algorithms fail when deployed incorrectly!

---

<sup>a</sup><https://appliedcryptography.page/paper/#imperfect-dh>



# Logjam finding: 512-bit primes still in use

2,631 servers with browser-trusted certificates used 512-bit primes!

- **Why this is bad:**
  - 512-bit discrete logarithm is solvable with modest resources
  - Passive eavesdropper can decrypt connections after the fact
  - No active attack needed—just record and decrypt later!
- **Real-world proof:** Researchers decrypted test connections to [www.fbi.gov](http://www.fbi.gov)
  - FBI website used OpenSSL's default 512-bit group until April 2015
  - Researchers had precomputed the discrete log table for this group
  - Complete compromise of "secure" government communications!
- **The lesson:** Default configurations can be dangerously weak
  - Many admins never change defaults
  - OpenSSL's 512-bit default was catastrophically inadequate

# Another Logjam finding: Composite-order subgroup attacks

- Not all servers use “safe” primes where  $(p - 1)/2$  is also prime
  - 4,800 out of 70,000 primes were not safe
  - If  $p - 1$  has small factors, the group has weak subgroups
  - **Note:** Even with non-safe primes, proper generator choice can maintain security
- Non-safe primes used with incorrectly chosen  $g$ :
  - Servers use short exponents (160-256 bits) for “efficiency”
  - **AND** generator  $g$  generates a small subgroup
  - If  $g$  generates the full group (order  $p - 1$ ), you’re still secure!
  - But many implementations chose  $g$  poorly
- **Real impact:** 159 servers fully compromised, including VPN devices and web conferencing servers!

# The dark side: unauthenticated Diffie-Hellman

Even with fully correct Diffie-Hellman implementations,  
there's a serious problem...

- **The vulnerability:** Basic DH has no authentication
  - Alice can't verify she's talking to Bob
  - Bob can't verify he's talking to Alice
- **The attack:** Man-in-the-middle (MITM)
  - Mallory sits between Alice and Bob
  - Alice does DH with Mallory, thinking it's Bob
  - Bob does DH with Mallory, thinking it's Alice
  - Mallory can read and modify everything!
- **Real-world impact:** This attack is practical and devastating!

# Man-in-the-middle attack on DH

How Mallory breaks “secure” communication:

1. **Alice → Mallory:** Alice sends  $g^a$  (thinking it goes to Bob)
2. **Mallory → Bob:** Mallory sends  $g^m$  (Bob thinks it's from Alice)
3. **Bob → Mallory:** Bob sends  $g^b$  (thinking it goes to Alice)
4. **Mallory → Alice:** Mallory sends  $g^m$  (Alice thinks it's from Bob)
5. **Result:**
  - Alice and Mallory share secret  $g^{am}$
  - Bob and Mallory share secret  $g^{bm}$
  - Alice and Bob don't share any secret!
6. **Communication:** Alice encrypts with  $g^{am}$ , Mallory decrypts, reads/modifies, re-encrypts with  $g^{bm}$  for Bob

**Alice and Bob never know they've been compromised!**

# Why MITM attacks succeed

- **Public values look random:**  $g^a$  and  $g^m$  are indistinguishable.
  - Both appear to be random group elements.
  - No way to tell if they come from the intended party.
- **No identity verification:** DH only establishes a shared secret.
  - Doesn't prove who you're sharing it with!
  - Like agreeing on a secret handshake with someone wearing a mask.
- **Active vs. passive attacks:**
  - DH protects against **passive** eavesdropping.
  - Does nothing against **active** manipulation.
- **Historical impact:** This attack has compromised real systems for decades.

# Solution: Authenticated Key Exchange

## Authenticated Key Exchange (AKE)

Key exchange that verifies the identity of the parties involved, preventing man-in-the-middle attacks.

- **Core idea:** Combine DH with authentication mechanisms
- **Common approaches:**
  - **Digital signatures:** Sign the DH public values (TLS).
  - **Pre-shared keys:** Use existing shared secrets (IPsec).
  - **Certificates:** Use a trusted third party (Certificate Authority in HTTPS).
  - **Password-based:** Derive authentication from passwords (SRP protocols).
- **Goal:** Ensure that Alice and Bob can verify they're really talking to each other.

# TLS: authenticated DH with certificates


## How HTTPS prevents MITM attacks:

1. **Server authentication:** Gmail sends its certificate along with  $g^b$ 
  - Certificate proves “this DH value really came from gmail.com”
  - Signed by a trusted Certificate Authority (CA)
2. **Certificate verification:** Your browser checks:
  - Is the signature valid?
  - Is the CA trusted?
  - Does the certificate match “gmail.com”?
  - Has the certificate expired?
3. **If verification passes:** You know you’re really talking to Gmail
4. **If verification fails:** Browser shows scary warnings!

**Result:** MITM attacks become much harder (but not impossible!)

# Signal: authenticated DH with fingerprints

- **The bootstrapping problem:** How do Alice and Bob initially authenticate?
  - No pre-existing certificates.
  - No trusted third parties.
- **Signal's solution:** Security numbers (fingerprints)
  - Each conversation gets a unique 60-digit number.
  - Derived from both parties' long-term identity keys.
- **Manual verification:** Users compare numbers out-of-band.
  - Read over the phone...
  - Show in person...
  - Send via different app...



Signal security number verification screen.



# SSH: authenticated DH with host keys

## How SSH prevents server impersonation:

- **First connection:** Server presents its “host key” along with DH public value
  - SSH shows you a fingerprint: SHA256:ABC123...
  - You’re supposed to verify this out-of-band (but nobody does!)
- **Trust on first use (TOFU):** Client remembers the host key
  - Stored in `~/.ssh/known_hosts`
- **Subsequent connections:** Client checks if host key matches
  - If different, gives you a heart attack: `WARNING: REMOTE HOST IDENTIFICATION HAS CHANGED!`
  - If same: Connection proceeds normally
- **User authentication:** Usually with passwords or public keys

**Weakness:** TOFU is vulnerable on the very first connection!

Section 2

Diffie-Hellman

---

Subsection 2.3

The Future of Diffie-Hellman

# Modern implementations: elliptic curves

## Traditional DH

- Uses  $\mathbb{Z}_p^*$  (integers mod  $p$ )
- Requires 2048+ bit numbers
- Slower computations
- Larger public keys

## Elliptic Curve DH (ECDH)

- Uses elliptic curve groups
- 256-bit keys  $\approx$  3072-bit traditional DH
- Much faster computations
- Smaller public keys, less bandwidth

- **Popular curves:** P-256, P-384, X25519, X448
- **Same security:** Based on elliptic curve discrete logarithm problem
- **Real-world adoption:** ECDH is now standard in TLS, Signal, etc.
- **Performance matters:** Especially important for mobile devices and IoT

# The quantum threat to Diffie-Hellman

## All DH variants are doomed...

- **Shor's algorithm** (1994) can break DH on quantum computers.
  - Solves discrete logarithm in polynomial time.
  - Works for both traditional DH and ECDH.
- **Timeline concerns:**
  - Large quantum computers don't exist yet.
  - But adversaries might store encrypted data now, decrypt later.
  - "Harvest now, decrypt later" attacks.
- **Post-quantum key exchange:** New algorithms under development.
  - ML-KEM (based on lattice problems)
  - SIDH/SIKE (based on isogenies, but recently broken!)
  - Code-based and hash-based alternatives

# Lessons from 50 years of Diffie-Hellman

- **Elegant mathematics:** Simple idea with profound implications.
  - Two numbers raised to secret powers in a mathematical group.
- **Security requires more than math:** Authentication is crucial.
  - Pure DH is vulnerable to active attacks.
  - Real systems need identity verification.
- **Efficiency drives adoption:** Elliptic curves made DH practical everywhere.
  - Performance improvements enable new applications.
- **Future challenges:** Quantum computers will force reinvention.
  - But the core insight—shared secrets from public exchanges—will survive.
- **Cryptography is a living field:** Continuous evolution and adaptation.

# From hard problems to real-world security

## The journey we've traced

1. **Mathematical insight:** Discrete logarithm is hard to compute.
2. **Cryptographic innovation:** Diffie-Hellman key exchange leverages this hardness.
3. **Real-world impact:** Secure communication for billions of people daily.

**This is the power of applied cryptography:** transforming abstract mathematical problems into tools that help people and protect our digital lives.

Section 3

RSA

# RSA: The other public-key pioneer

## One year after Diffie-Hellman came RSA (1977)

- **Named after:** Rivest, Shamir, and Adleman (MIT researchers)
- **Different approach:** Based on the hardness of factoring large numbers
- **Key innovation:** Combines public-key encryption AND digital signatures
- **Revolutionary impact:** Enabled secure communication without prior key exchange

### RSA's Core Idea

Anyone can encrypt messages to you using your public key, but only you can decrypt them with your private key.



# How RSA works (simplified)

## The mathematical foundation:

### 1. Key generation:

- Choose two large primes  $p$  and  $q$
- Compute  $n = p \times q$  (public modulus)
- Choose public exponent  $e$  (commonly 65537)
- Compute private exponent  $d$  such that  $e \times d \equiv 1 \pmod{\phi(n)}$ 
  - $\phi(n)$  is Euler's totient function: counts integers less than  $n$  that are coprime to  $n$
  - For  $n = p \times q$ :  $\phi(n) = (p - 1)(q - 1)$

2. **Encryption:**  $c = m^e \bmod n$  (anyone can do this)

3. **Decryption:**  $m = c^d \bmod n$  (only private key holder can do this)

**Security assumption:** Factoring  $n$  back into  $p$  and  $q$  is computationally infeasible

# RSA's dual functionality

## Encryption

- Public key encrypts
- Private key decrypts
- Provides confidentiality
- Anyone can send you secrets

## Digital Signatures

- Private key signs
- Public key verifies
- Provides authenticity
- You can prove you wrote something

- **Historical significance:** First algorithm to do both!
- **Widespread adoption:** Used in SSL/TLS, email encryption, code signing
- **Patent controversy:** Was patented until 2000, limiting early adoption

# RSA in practice

## Where RSA dominated for decades:

- **HTTPS certificates:** Most websites used RSA signatures until recently
- **Email encryption:** PGP/GPG relied heavily on RSA
- **Software signing:** Windows, macOS, Linux packages
- **SSH keys:** RSA was the default for years

## Key sizes and security:

- 1024-bit RSA: Deprecated, considered insecure since 2010
- 2048-bit RSA: Current minimum, but showing its age
- 4096-bit RSA: Secure but slow and unwieldy
- Compare to ECDSA: 256-bit keys offer similar security to 3072-bit RSA!

# Why we're not covering RSA in depth

## RSA is a cryptographic dinosaur

- **On its way out:** Being actively replaced everywhere
  - TLS 1.3 prefers ECDSA and EdDSA
  - Modern systems default to elliptic curves
  - Even RSA Inc. recommends moving away from RSA!
- **Not a building block:** Unlike Diffie-Hellman
  - DH principles apply to modern protocols (ECDH, X25519)
  - RSA's approach doesn't translate to newer systems
- **Performance issues:** Terribly slow compared to alternatives
  - 4096-bit RSA signature: ~10ms
  - Ed25519 signature: ~0.1ms (100x faster!)
- **Implementation pitfalls:** Padding oracle attacks, timing attacks, etc.
- **The future is elliptic curves and post-quantum:** Time better spent there!



# Applied Cryptography

Summer 2026

Part 1: Provable Security

## 1.7: Hard Problems & Diffie-Hellman

Nadim Kobeissi

<https://appliedcryptography.page>